


# Joining data frames



# Housekeeping

- Wrangling coding practice due tonight!
- Office hours tomorrow 10am-12pm
- Problem Set 3 has large R component

# Joining data frames

Assume we have two data frame, `x` and `y`. There are some shared variables (i.e. columns) in the two. Suppose we want to combine them together into one single data frame.

```
1 something_join(x, y)
```

# Setup

For the next few slides...

1	x
---	---

ID	x_val
1	x1
2	x2
3	x3

1	y
---	---

ID	y_val
1	y1
2	y2
4	y4

# left\_join()

Adds columns to **x** from **y**, matching all rows in **x**

```
1 x
```

ID	x_val
1	x1
2	x2
3	x3

```
1 y
```

ID	y_val
1	y1
2	y2
4	y4

```
1 left_join(x, y, by = "ID")
```

ID	x_val	y_val
1	x1	y1
2	x2	y2
3	x3	NA

left\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

# right\_join()

Adds columns to **x** from **y**, matching all rows in **y**

```
1 x
```

ID	x_val
1	x1
2	x2
3	x3

```
1 y
```

ID	y_val
1	y1
2	y2
4	y4

```
1 right_join(x, y, by = "ID")
```

ID	x_val	y_val
1	x1	y1
2	x2	y2
4	NA	y4

right\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

# full\_join()

Adds columns to **x** from **y**, matching all rows in **x** OR **y**

```
1 x
```

ID	x_val
1	x1
2	x2
3	x3

```
1 y
```

ID	y_val
1	y1
2	y2
4	y4

```
1 full_join(x, y, by = "ID")
```

ID	x_val	y_val
1	x1	y1
2	x2	y2
3	x3	NA
4	NA	y4

full\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

# inner\_join()

All rows from **x** where there are matching values in **y**, return all combination of multiple matches in the case of multiple matches

```
1 x
```

ID	x_val
1	x1
2	x2
3	x3

```
1 y
```

ID	y_val
1	y1
2	y2
4	y4

```
1 inner_join(x, y, by = "ID")
```

ID	x_val	y_val
1	x1	y1
2	x2	y2

inner\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4



# inner\_join() (cont.)

Example with multiple matches:

```
1 x2
```

ID	x_val
1	x1
2	x2
3	x3
1	new_x

```
1 y2
```

ID	y_val
1	y1
2	y2
4	y4
1	new_y

```
1 inner_join(x2, y2, by = "ID")
```

Warning in inner\_join(x2, y2, by = "ID"): Detected an unexpected many-to-many relationship between `x` and `y`.

i Row 1 of `x` matches multiple rows in `y`.

i Row 1 of `y` matches multiple rows in `x`.

i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to silence this warning.

ID	x_val	y_val
1	x1	y1
1	x1	new_y
2	x2	y2
1	new_x	y1
1	new_x	new_y

# semi\_join()

Returns all rows from **x** with a match in **y**, but does not add columns from **y**

```
1 x
```

ID	x_val
1	x1
2	x2
3	x3

```
1 y
```

ID	y_val
1	y1
2	y2
4	y4

```
1 semi_join(x, y, by = "ID")
```

ID	x_val
1	x1
2	x2

semi\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

# anti\_join()

Returns all rows from **x** **without** any match in **y**, and not add columns from **y**

```
1 x
```

ID	x_val
1	x1
2	x2
3	x3

```
1 y
```

ID	y_val
1	y1
2	y2
4	y4

```
1 anti_join(x, y, by = "ID")
```

ID	x_val
3	x3

anti\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

# Joining with different variable names

If the variables in **x** and **y** have different names but we know they represent the same variable:

```
1 x
```

ID	x_val
1	x1
2	x2
3	x3

```
1 y3
```

ID_y	y_val
1	y1
2	y2
4	y4

```
1 left_join(x, y3, by = c("ID" = "ID_y"))
```

ID	x_val	y_val
1	x1	y1
2	x2	y2
3	x3	NA

# Joining on multiple variables

Can specify more than one variable in the `by` argument. Will need to a vector of character objects.

```
1 enrollment
```

student_id	course	start_year
1	STAT 310	F22
1	MATH 223	F22
2	STAT 310	F23
3	STAT 201	W24

```
1 payment
```

student_id	Course	status
1	STAT 310	paid
1	MATH 223	paid
2	STAT 310	unpaid
3	STAT 201	paid

```
1 # note, multiple join functions would work in this example!  
2 inner_join(enrollment, payment, by = c("student_id", "course" = "Course"))
```

student_id	course	start_year	status
1	STAT 310	F22	paid
1	MATH 223	F22	paid
2	STAT 310	F23	unpaid
3	STAT 201	W24	paid

**Live code**

# Example

We have data on fishery harvests (in tons) by countries from 2016:

```
1 fish |>  
2   slice(1:9)
```

country	capture	aquaculture
Afghanistan	1000	1200
Albania	7886	950
Algeria	95000	1361
American Samoa	3047	20
Andorra	0	0
Angola	486490	655
Antigua and Barbuda	3000	10
Argentina	755226	3673
Armenia	3758	16381

# Bringing in continent

Suppose I would like to explore the data on a continent level. We don't have continent in the current data frame, but we could join in the following data:

```
1 continents |>  
2   slice(1:5)
```

country	continent
Afghanistan	Asia
Åland Islands	Europe
Albania	Europe
Algeria	Africa
American Samoa	Oceania

- We want to keep all rows and columns from `fish` and add a column for corresponding continents. Which join function should we use?
- We want to keep all rows from `fish` for which we have a corresponding continent and add a column for corresponding continents. Which join function should we use?



# Example (cont.)

```
1 left_join(fish, continents, by = "country") |>
2   slice(1:9)
```

country	capture	aquaculture	continent
Afghanistan	1000	1200	Asia
Albania	7886	950	Europe
Algeria	95000	1361	Africa
American Samoa	3047	20	Oceania
Andorra	0	0	Europe
Angola	486490	655	Africa
Antigua and Barbuda	3000	10	NA
Argentina	755226	3673	Americas
Armenia	3758	16381	Asia

```
1 inner_join(fish, continents, by = "country") |>
2   slice(1:9)
```

country	capture	aquaculture	continent
Afghanistan	1000	1200	Asia
Albania	7886	950	Europe
Algeria	95000	1361	Africa
American Samoa	3047	20	Oceania
Andorra	0	0	Europe
Angola	486490	655	Africa
Argentina	755226	3673	Americas
Armenia	3758	16381	Asia
Aruba	142	0	Americas

- Notice the **NA**
- Could also use the following piping code:

```
1 fish |>
2   left_join(continents, by = "country")
```


# group\_by()

```
1 fish |>
2   left_join(continents, by = "country") |>
3   group_by(continent) |> # get continent-level summary stats
4   summarise(mean_capture = mean(capture), sd_capture = sd(capture))
```

```
# A tibble: 6 × 3
  continent mean_capture sd_capture
  <chr>         <dbl>         <dbl>
1 Africa      180705.      266107.
2 Americas    433235.     1038899.
3 Asia       1036018.     2869652.
4 Europe      317874.      797072.
5 Oceania      74660.      121027.
6 <NA>       134722.      448079.
```

# Visualize

```
1 fish |>  
2   na.omit() |> #remove observations with any NAs  
3   left_join(continent, by = "country") |>  
4   ggplot(aes(x = continent, y = capture)) +  
5   geom_boxplot(outliers = F) +  
6   theme_bw() +  
7   labs(caption = "Excluding outliers")
```




Excluding outliers

# Visualize (cont.)

```
1 fish |>
2   mutate(log_capture = log(capture)) |>
3   left_join(continents, by = "country") |>
4   na.omit() |> #remove observations with any NAs
5   ggplot(aes(x = continent, y = capture)) +
6   geom_boxplot(outliers = F) +
7   theme_bw() +
8   labs(caption = "Excluding outliers", y = "Log capture")
```

Warning: Removed 4 rows containing non-finite outside the scale range (`stat\_boxplot()`).



Excluding countries with 0 capture

